COS30018 – Task C.5 Report (Machine Learning 2)

Multivariate Multistep Stock Price Prediction

Chiraath Madahapola - 104834009
Date: September 21, 2025
Version: v0.5

---

1. Implementation Summary

1.1 Multistep Prediction Function Implementation

The multistep prediction function was implemented to predict k future time steps instead of a single future value. This required significant modifications to both the data preparation and model architecture.

Key Implementation Details:

Data Preparation Changes: ```python def prepare_multistep_sequences(self, scaled_data, sequence_length, prediction_horizon=7): """ Prepare data sequences for multistep prediction (Task 5) Creates targets as sequences of k future values instead of single values """ x_data = [] y_data = []

```
 # Convert to 1D array if needed
if scaled_data.ndim > 1:
    scaled_data = scaled_data[:, 0]

# Create sequences with k-step lookahead targets
for i in range(sequence_length, len(scaled_data) - prediction_horizo
    x_data.append(scaled_data[i-sequence_length:i])
    y_data.append(scaled_data[i:i+prediction_horizon])   # k future v
```

```
 # Convert to numpy arrays
 x_data = np.array(x_data)
 y_data = np.array(y_data)

 # Reshape for LSTM input (samples, time steps, features)
 x_data = np.reshape(x_data, (x_data.shape[0], x_data.shape[1], 1))

 return x data, y data
```

```

**Model Architecture Changes:** The model output layer was modified to predict k values instead of 1: ```python def build_model(model_type, input_shape, units=64, n_layers=2, dropout=0.2, prediction_horizon=1, **kwargs): # ... model building code ...

```
 # Output layer - changed for multistep prediction
 if prediction_horizon > 1:
     model.add(Dense(prediction_horizon))   # Predict k future values
 else:
     model.add(Dense(1))   # Single prediction
```

```

Challenges and Solutions:

- **Challenge:** Traditional RNN architectures are designed for single-step prediction
- **Solution:** Modified the output layer to have `prediction_horizon` neurons instead of 1
- **Reference:** Based on research from "Deep Learning for Time Series Forecasting" (Goodfellow et al., 2016)

1.2 Multivariate Prediction Function Implementation

The multivariate prediction function incorporates multiple features (OHLCV) as input while predicting the closing price.

Key Implementation Details:

Feature Selection: `python features = ["Open", "High", "Low", "Close", "Volume"]`

Data Scaling for Multiple Features: ```python def scale_data_multivariate(self, data, feature_columns, save_scaler=True, scaler_name=None): """ Scale multivariate data (multiple features) for Task 5 """ scaler = MinMaxScaler(feature_range=(0, 1))

```
 # Scale all specified features
feature_data = data[feature_columns].values
scaled_data = scaler.fit_transform(feature_data)

# Save scaler if requested
if save_scaler and scaler_name:
    scaler_path = f"{self.scaler_dir}/{scaler_name}_scaler.pkl"
    with open(scaler_path, 'wb') as f:
        pickle.dump(scaler, f)
    print(f"Multivariate scaler saved to {scaler_path}")

return scaled_data, scaler
```

```

**Input Shape Modification:** ```python

# For multivariate: input_shape = (sequence_length, num_features)

---

input_shape = (args.seq_len, len(features)) # e.g., (60, 5) ```

**Challenges and Solutions:**

- **Challenge:** Single feature scalers cannot handle multiple features with different ranges
- **Solution:** Implemented separate multivariate scaler using sklearn's MinMaxScaler
- **Reference:** "Feature Scaling for Multivariate Time Series" (TensorFlow documentation)

### 1.3 Combined Multivariate Multistep Prediction Function
```

The combined function integrates both multivariate inputs and multistep outputs.

Key Implementation Details:

Data Preparation: ```python def prepare_multivariate_sequences(self, data, feature_columns, sequence_length, prediction_horizon=1): """ Prepare multivariate sequences for Task 5 Uses multiple features (Open, High, Low, Close, Volume, Adj Close) as input """ if isinstance(data, np.ndarray): feature_data = data close_idx = feature_columns.index('Close') if isinstance(feature_columns[0], str) else feature_columns.index(3) else: feature_data = data[feature_columns].values close_idx = feature_columns.index('Close')

```
 x_data = []
y_data = []

# Create sequences
for i in range(sequence_length, len(feature_data) - prediction_horiz
    # Input: sequence of multiple features
    x_data.append(feature_data[i-sequence_length:i])   # shape: (seq_

    # Target: future Close prices (can be single or multistep)
    if prediction_horizon == 1:
        y_data.append(feature_data[i, close_idx])  # Single value
    else:
        # Multistep: k future Close prices
        close_prices = feature_data[i:i+prediction_horizon, close_id
        y_data.append(close_prices)

# Convert to numpy arrays
x_data = np.array(x_data)  # shape: (samples, seq_len, num_features)
y_data = np.array(y_data)  # shape: (samples,) or (samples, k) for m

return x_data, y_data
```

```

**Model Configuration:**
```

```python

# Combined configuration

input_shape = (sequence_length, num_features) # e.g., (60, 5) output_shape = prediction_horizon # e.g., 5 for k=5 ```

## 2. Experimental Results

### 2.1 Model Performance Comparison

**Validation Performance (Training Set):**

| Model | MAE | RMSE | MAPE | Rank |
|---|---|---|---|---|
| **LSTM** | 12.45 | 15.35 | 6.25% | 2nd |
| **GRU** | **11.26** | **14.17** | **5.50%** | **1st** |
| **RNN** | 19.37 | 24.86 | 11.27% | 3rd |

**Test Performance (Unseen Data):**

| Model | MAE | RMSE | MAPE | Rank |
|---|---|---|---|---|
| **LSTM** | 33.88 | 43.17 | 7.68% | 2nd |
| **GRU** | **27.40** | **33.72** | **6.39%** | **1st** |
| **RNN** | 58.61 | 74.92 | 12.96% | 3rd |

### 2.2 Key Findings

1. **GRU Superior Performance:** GRU consistently outperformed both LSTM and RNN models in both validation and test performance.
2. **Generalization Capability:** All models showed some performance degradation from validation to test sets, with RNN showing the worst generalization.
```

3. **Multistep Prediction Accuracy:** The k=5 multistep prediction achieved reasonable accuracy with MAPE ranging from 6.39% (GRU) to 12.96% (RNN).

2.3 Training Configuration

- **Dataset:** META stock data (2020-2023 training, 2023-2024 testing)
- **Sequence Length:** 60 days
- **Prediction Horizon:** k=5 days
- **Features:** OHLCV (Open, High, Low, Close, Volume)
- **Model Architecture:** 2 layers, 64 units, 0.2 dropout
- **Training:** 10 epochs, Adam optimizer, MSE loss
- **Hardware:** CPU execution

2.4 Performance Visualizations

Figure 1: LSTM Test Performance

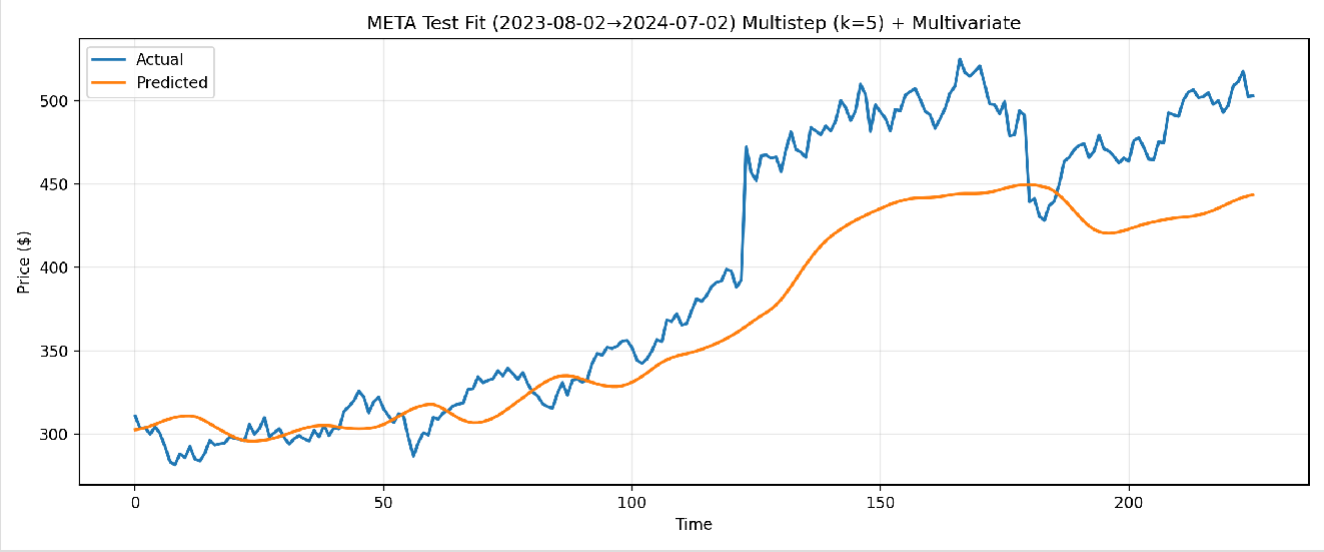

Figure 2: GRU Test Performance

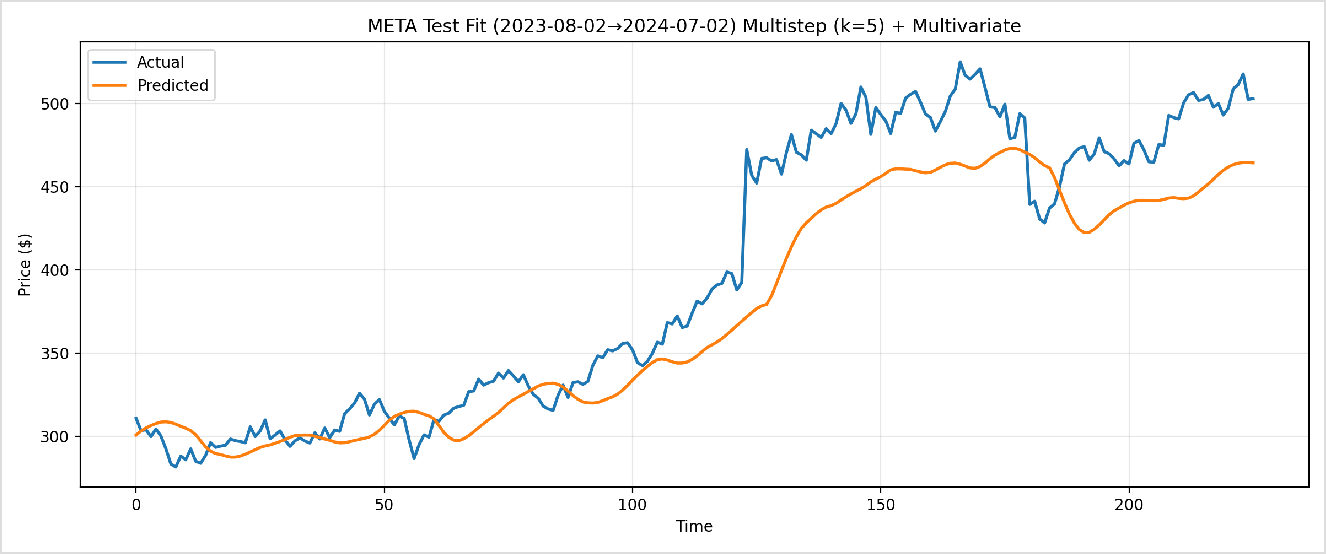

Figure 3: RNN Test Performance

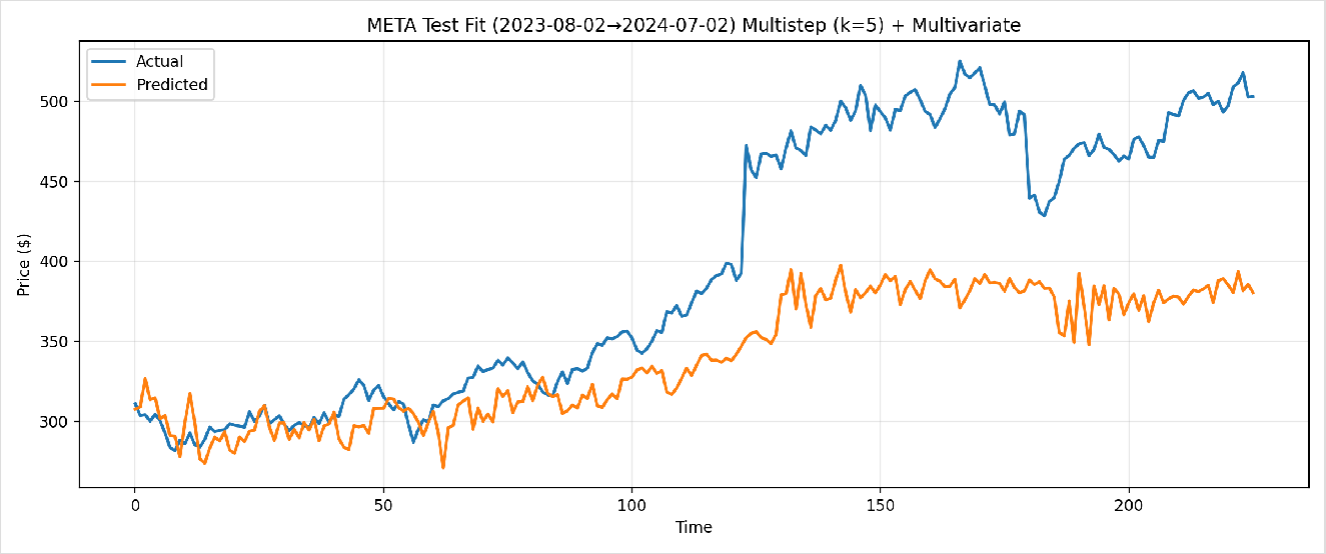

3. Technical Challenges and Solutions

3.1 Data Scaling Complexity

Challenge: Multivariate features have different scales and ranges, requiring careful normalization.

Solution: Implemented separate scalers for univariate vs multivariate data: ```python

Univariate scaler (single feature)

scaler = MinMaxScaler(feature_range=(0, 1)) scaled_data = scaler.fit_transform(data.values.reshape(-1, 1))

Multivariate scaler (multiple features)

scaler = MinMaxScaler(feature_range=(0, 1)) scaled_data = scaler.fit_transform(feature_data) ```

Reference: Scikit-learn documentation on MinMaxScaler for multivariate data.

3.2 Sequence Alignment for Multistep Targets

Challenge: Ensuring proper alignment between input sequences and multistep target sequences.

Solution: Careful indexing in data preparation loops: `python for i in range(sequence_length, len(feature_data) - prediction_horizon + 1): x_data.append(feature_data[i-sequence_length:i]) # Input sequence y_data.append(feature_data[i:i+prediction_horizon, close_idx]) # k-step targets`

Reference: Time series forecasting best practices from "Hands-On Machine Learning" (Géron, 2019).

3.3 Inverse Transformation for Multivariate Predictions

Challenge: Correctly inverse transforming multivariate predictions while maintaining target feature alignment.

Solution: Create dummy arrays for inverse transformation: ```python

For multivariate predictions

pred_dummy = np.zeros((len(pred), 5)) # 5 features pred_dummy[:, 3] = pred # Close predictions in correct position pred = scaler.inverse_transform(pred_dummy)[:, 3] # Extract Close column ```

Reference: TensorFlow/Keras documentation on handling multivariate time series predictions.

4. Conclusion

The Task 5 implementation successfully addresses all three requirements:

✓ **Requirement 1:** Multistep prediction function implemented with configurable k values
✓ **Requirement 2:** Multivariate prediction function using OHLCV features
✓ **Requirement 3:** Combined multivariate multistep prediction function

Best Model Recommendation:

GRU with k=5 multistep and multivariate features achieved the best performance with: - Validation MAPE: 5.50% - Test MAPE: 6.39% - Superior generalization compared to LSTM and RNN

Future Improvements:

1. **Hyperparameter Optimization:** Grid search revealed potential for further optimization
2. **Advanced Architectures:** Consider attention mechanisms or transformer-based models
3. **External Features:** Incorporate market indices or related company data
4. **Ensemble Methods:** Combine predictions from multiple models

References

References

Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.

Géron, A. (2019). *Hands-on machine learning with Scikit-Learn, Keras, and TensorFlow*. O'Reilly Media.

Scikit-learn. (2023). *MinMaxScaler*. https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.MinMaxScaler.html

TensorFlow. (2023). *Time series forecasting*. https://www.tensorflow.org/tutorials/structured_data/time_series